



 [ovh](#) / [ai-training-examples](#) Public


<> Code


 Issues


 Pull requests 1

 Actions

 Security

 Insights

[ai-training-examples](#) / [notebooks](#) / [natural-language-processing](#) / [llm](#) / [miniconda](#) / [llama2-fine-tuning](#) / [llama_2_finetuning.ipynb](#) 



 **MathieuBsqt** last month





2.13 MB



Fine-Tuning LLaMA 2 Models with a single GPU and OVHcloud

In this OVHcloud tutorial, we will walk you through the process of fine-tuning LLaMA 2 models, providing step-by-step instructions.



Introduction

On July 18, 2023, [Meta](#) released [LLaMA 2](#), the latest version of their Large Language Model (LLM).

Trained between January 2023 and July 2023 on 2 trillion tokens, these new models outperforms other LLMs on many benchmarks, including reasoning, coding, proficiency, and knowledge tests. This release comes in different flavors, with parameter sizes of 7B, 13B, and a mind-blowing 70B. Models are intended for free for both commercial and research use in English.

To suit every text generation needed and fine-tune these models, we will use [QLoRA \(Efficient Finetuning of Quantized LLMs\)](#), a highly efficient fine-tuning technique that involves quantizing a pretrained LLM to just 4 bits and adding small "Low-Rank Adapters". This unique approach allows for fine-tuning LLMs using just a single GPU! This technique is supported by the [PEFT library](#).

Requirements

To successfully fine-tune LLaMA 2 models, you will need the following:

- **Set up your Python environment** by installing the `requirements.txt` file
- **Llama 2 Model.** To obtain the Llama 2 model, you will need to:
 - Fill Meta's form to [request access to the next version of Llama](#). Indeed, the use of Llama 2 is governed by the Meta license, that you must accept in order to download the model weights and tokenizer.
 - Have a [Hugging Face](#) account (with the same email address you entered in Meta's form).
 - Have a [Hugging Face token](#).
 - Visit the page of one of the LLaMA 2 available models (version [7B](#), [13B](#) or [70B](#)) and accept Hugging Face's license terms and acceptable use

Once you have accepted this, you will get the following message: *Your request to access this repo has been successfully submitted, and is pending a review from the repo's authors*, which a few hours later should change to: *You have been granted access to this model*.

- Log in to the Hugging Face model Hub from your notebook's terminal. To do this, just click the `+` button and open a terminal. You can also perform this by clicking `File > New > Terminal`. Then, use the `huggingface-cli login` command, and enter your token. You will not need to add your token as git credential.

- **Powerful Computing Resources:** Fine-tuning the Llama 2 model requires substantial computational power. Ensure you are running code on GPU(s).

```
In [ ]: # Set up Python environment
!pip install -r requirements.txt
```

```
In [2]: # Import libraries
import argparse
import bitsandbytes as bnb
from datasets import load_dataset
from functools import partial
import os
from peft import LoraConfig, get_peft_model, prepare_model_for_kbit_training
import torch
from transformers import AutoModelForCausalLM, AutoTokenizer, set_seed, DataCollatorForLanguageModeling, Trainer, TrainingArguments
from datasets import load_dataset

# Reproducibility
seed = 42
set_seed(seed)
```

Download LLaMA 2 model

As mentioned before, LLaMA 2 models come in different flavors which are 7B, 13B, and 70B. Your choice can be influenced by your computational resources. Indeed, larger models require more resources, memory, processing power, and training time.

To download the model you have been granted access to, **make sure you are logged in to the Hugging Face model hub**. As mentioned in the requirements step, you need to use the `huggingface-cli login` command.

The following function will help us to download the model and its tokenizer. It requires a `bitsandbytes` configuration that we will define later

```
In [3]: def load_model(model_name, bnb_config):
n_gpus = torch.cuda.device_count()
max_memory = f'{40960}MB'

model = AutoModelForCausalLM.from_pretrained(
    model_name,
    quantization_config=bnb_config,
    device_map="auto", # dispatch efficiently the model on the
    max_memory = {i: max_memory for i in range(n_gpus)},
)
tokenizer = AutoTokenizer.from_pretrained(model_name, use_auth_

# Needed for LLaMA tokenizer
tokenizer.pad_token = tokenizer.eos_token

return model, tokenizer
```

Download a Dataset

There are many datasets that can help you fine-tune your model. You can even use your own dataset!

In this tutorial, we are going to download and use the [Databricks Dolly 15k dataset](#), which contains 15,000 prompt/response pairs. It was crafted by over 5,000 [Databricks](#) employees during March and April of 2023.

This dataset is designed specifically for fine-tuning large language models. Released under the [CC BY-SA 3.0 license](#), it can be used, modified, and extended by any individual or company, even for commercial applications. So it's a perfect fit for our use case!

However, like most datasets, this one has its limitations. Indeed, pay attention to the following points:

- It consists of content collected from the public internet, which means it may contain objectionable, incorrect or biased content and typo errors, which could influence the behavior of models fine-tuned using this dataset.
- Since the dataset has been created for Databricks by their own employees, it's worth noting that the dataset reflects the interests and semantic choices of Databricks employees, which may not be representative of the global population at large.
- We only have access to the `train` split of the dataset, which is its largest subset.

```
In [4]: # Load the databricks dataset from Hugging Face
from datasets import load_dataset

dataset = load_dataset("databricks/databricks-dolly-15k", split="tr
```

```
Downloading readme: 0%|          | 0.00/8.20k [00:00<?, ?B/s]
Downloading and preparing dataset json/databricks--databricks-dolly-1
5k to /workspace/.cache/huggingface/datasets/databricks__json/databr
icks_databricks_dolly_15k_7427226057624282/0_0_0/8bb11242116d547c741
```

```
LSKS--databricks-dolly-15k-7427aa6e57c34282/0.0.0/8bb11242116d547c741b2e8a1f18598ffdd40a1d4f2a2872c7a28b697434bc96...
Downloading data files: 0%|          | 0/1 [00:00<?, ?it/s]
Downloading data: 0%|          | 0.00/13.1M [00:00<?, ?B/s]
Extracting data files: 0%|          | 0/1 [00:00<?, ?it/s]
Generating train split: 0 examples [00:00, ? examples/s]
Dataset json downloaded and prepared to /workspace/.cache/huggingface/datasets/databricks__json/databricks--databricks-dolly-15k-7427aa6e57c34282/0.0.0/8bb11242116d547c741b2e8a1f18598ffdd40a1d4f2a2872c7a28b697434bc96. Subsequent calls will reuse this data.
```

Explore dataset

Once the dataset is downloaded, we can take a look at it to understand what it contains:

```
In [5]: print(f'Number of prompts: {len(dataset)}')
        print(f'Column names are: {dataset.column_names}')
```

Number of prompts: 15011
Column names are: ['instruction', 'context', 'response', 'category']

As we can see, it is composed of 4 fields named `instruction` , `context` , `response` , `category` . Let's take a look at 3 samples to better understand what we are talking about:

```
In [6]: import random
import pandas as pd

# Generate random indices
nb_samples = 3
random_indices = random.sample(range(len(dataset)), nb_samples)
samples = []

for idx in random_indices:
    sample = dataset[idx]

    sample_data = {
        'instruction': sample['instruction'],
        'context': sample['context'],
        'response': sample['response'],
        'category': sample['category']
    }
    samples.append(sample_data)

# Create a DataFrame and display it
df = pd.DataFrame(samples)
display(df)
```

	instruction	context	response	category
0	What athlete created the 'beast quake' for the...		Marshan Lynch	open_qa
1	Who wrote Democracy in America?		Alexis de Tocqueville wrote Democracy in America	open_qa
2	What links Brazil, Uruguay, Mozambique and Angola		Colonies of Portugal	open_qa

As we can see, each sample is a dictionary that contains:

- **An instruction:** What could be entered by the user, such as a question
- **A context:** Help to interpret the sample
- **A response:** Answer to the instruction
- **A category:** Classify the sample between Open Q&A, Closed Q&A, Extract information from Wikipedia, Summarize information from Wikipedia, Brainstorming, Classification, Creative writing

In [7]:

```
print(sample)
```

```
{'instruction': 'What links Brazil, Uruguay, Mozambique and Angola',
'context': '', 'response': 'Colonies of Portugal', 'category': 'open_
qa'}
```

If you look at several samples, you will see that most do not contain any context . But don't worry, it doesn't matter. What we need to do now is to pre-process our data.

Pre-processing dataset

Instruction fine-tuning is a common technique used to fine-tune a base LLM for a specific downstream use-case.

It will help us to format our prompts.

In [8]:

```
def create_prompt_formats(sample):
    """
    Format various fields of the sample ('instruction', 'context',
    Then concatenate them using two newline characters
    :param sample: Sample dictionary
    """
    INTRO_BLURB = "Below is an instruction that describes a task. W
    INSTRUCTION_KEY = "### Instruction:"
    INPUT_KEY = "### Input:"
    RESPONSE_KEY = "### Response:"
    END_KEY = "### End"

    blurb = f"{INTRO_BLURB}"
    instruction = f"{INSTRUCTION_KEY}\n{sample['instruction']}"
    input_context = f"{INPUT_KEY}\n{sample['context']}" if sample["
    response = f"{RESPONSE_KEY}\n{sample['response']}"
    end = f"{END_KEY}"

    parts = [part for part in [blurb, instruction, input_context, r
    formatted_prompt = "\n\n".join(parts)

    sample["text"] = formatted_prompt

    return sample

print(create_prompt_formats(sample)["text"])
```

Below is an instruction that describes a task. Write a response that appropriately completes the request.

```

### Instruction:
What links Brazil, Uruguay, Mozambique and Angola

### Response:
Colonies of Portugal

### End

```

As we can see, each part is now delimited by hashtags that describe the prompt.

Now, we will use our model tokenizer to process these prompts into tokenized ones. The goal is to create input sequences of uniform length (which are suitable for fine-tuning the language model because it maximizes efficiency and minimize computational overhead), that must not exceed the model's maximum token limit.

```

In [9]: # SOURCE https://github.com/databricks/dolly/blob/master/traini
def get_max_length(model):
    conf = model.config
    max_length = None
    for length_setting in ["n_positions", "max_position_embeddings":
        max_length = getattr(model.config, length_setting, None)
        if max_length:
            print(f"Found max length: {max_length}")
            break
    if not max_length:
        max_length = 1024
        print(f"Using default max length: {max_length}")
    return max_length

def preprocess_batch(batch, tokenizer, max_length):
    """
    Tokenizing a batch
    """
    return tokenizer(
        batch["text"],
        max_length=max_length,
        truncation=True,
    )

# SOURCE https://github.com/databricks/dolly/blob/master/traini
def preprocess_dataset(tokenizer: AutoTokenizer, max_length: int, s
    """Format & tokenize it so it is ready for training
    :param tokenizer (AutoTokenizer): Model Tokenizer
    :param max_length (int): Maximum number of tokens to emit from
    """

    # Add prompt to each sample
    print("Preprocessing dataset...")
    dataset = dataset.map(create_prompt_formats)#, batched=True)

    # Apply preprocessing to each batch of the dataset & and remove
    _preprocessing_function = partial(preprocess_batch, max_length=
    dataset = dataset.map(
        _preprocessing_function,
        batched=True,
        remove_columns=["instruction", "context", "response", "text"
    )

    # Filter out samples that have input_ids exceeding max_length
    dataset = dataset.filter(lambda sample: len(sample["input_ids"])

```



```
dataset = dataset.filter(lambda sample: len(sample["input_ids"]) > 1)

# Shuffle dataset
dataset = dataset.shuffle(seed=seed)

return dataset
```

With these functions, our dataset will be ready for fine-tuning !

Create bnb config

This will allow us to load our LLM in 4 bits. This way, we can divide the used memory by 4 and import the model on smaller devices. We choose to apply bfloat16 compute data type and nested quantization for memory-saving purposes.

```
In [10]: def create_bnb_config():
        bnb_config = BitsAndBytesConfig(
            load_in_4bit=True,
            bnb_4bit_use_double_quant=True,
            bnb_4bit_quant_type="nf4",
            bnb_4bit_compute_dtype=torch.bfloat16,
        )

        return bnb_config
```

To leverage the LoRa method, we need to wrap the model as a PeftModel.

To do this, we need to implement a [LoRa configuration](#):

```
In [11]: def create_peft_config(modules):
        """
        Create Parameter-Efficient Fine-Tuning config for your model
        :param modules: Names of the modules to apply Lora to
        """
        config = LoraConfig(
            r=16, # dimension of the updated matrices
            lora_alpha=64, # parameter for scaling
            target_modules=modules,
            lora_dropout=0.1, # dropout probability for layers
            bias="none",
            task_type="CAUSAL_LM",
        )

        return config
```

Previous function needs the target modules to update the necessary matrices.

The following function will get them for our model:

```
In [12]: # SOURCE https://github.com/artidoro/qlora/blob/main/qlora.py
def find_all_linear_names(model):
    cls = bnb.nn.Linear4bit #if args.bits == 4 else (bnb.nn.Linear8
    lora_module_names = set()
    for name, module in model.named_modules():
        if isinstance(module, cls):
            names = name.split('.')
            lora_module_names.add(names[0] if len(names) == 1 else
```

```

if 'lm_head' in lora_module_names: # needed for 16-bit
    lora_module_names.remove('lm_head')
return list(lora_module_names)

```

Once everything is set up and the base model is prepared, we can use the `print_trainable_parameters()` helper function to see how many trainable parameters are in the model. We expect the `lora_model` to have fewer trainable parameters compared to the original one, since we want to perform fine-tuning.

```

In [13]: def print_trainable_parameters(model, use_4bit=False):
        """
        Prints the number of trainable parameters in the model.
        """
        trainable_params = 0
        all_param = 0
        for _, param in model.named_parameters():
            num_params = param.numel()
            # if using DS Zero 3 and the weights are initialized empty
            if num_params == 0 and hasattr(param, "ds_numel"):
                num_params = param.ds_numel

            all_param += num_params
            if param.requires_grad:
                trainable_params += num_params
        if use_4bit:
            trainable_params /= 2
        print(
            f"all params: {all_param:,d} || trainable params: {trainabl
)

```

Training

Now that everything is ready, we can pre-process our dataset and load our model using the set configurations.

Then, we can run our fine-tuning process.

```

In [15]: # Load model from HF with user's token and with bitsandbytes config
model_name = "meta-llama/Llama-2-7b-hf"
bnb_config = create_bnb_config()
model, tokenizer = load_model(model_name, bnb_config)

```

```

Downloading (...)lve/main/config.json: 0%|          | 0.00/609 [00:00
<?, ?B/s]
Downloading (...)fetensors.index.json: 0%|          | 0.00/26.8k [00:
00<?, ?B/s]
Downloading shards: 0%|          | 0/2 [00:00<?, ?it/s]
Downloading (...)of-000002.safetensors: 0%|          | 0.00/9.98G [00:
00<?, ?B/s]
Downloading (...)of-000002.safetensors: 0%|          | 0.00/3.50G [00:
00<?, ?B/s]
Loading checkpoint shards: 0%|          | 0/2 [00:00<?, ?it/s]
Downloading (...)neration_config.json: 0%|          | 0.00/167 [00:00
<?, ?B/s]

```

```

/workspace/.miniconda3/lib/python3.9/site-packages/transformers/model
s/auto/tokenization_auto.py:628: FutureWarning: The `use_auth_token`
argument is deprecated and will be removed in v5 of Transformers.

```

```
warnings.warn(
Downloading (...)okenizer_config.json: 0%|          | 0.00/776 [00:00
<?, ?B/s]
Downloading tokenizer.model: 0%|          | 0.00/500k [00:00<?, ?B/
s]
Downloading (...)main/tokenizer.json: 0%|          | 0.00/1.84M [00:
00<?, ?B/s]
Downloading (...)cial_tokens_map.json: 0%|          | 0.00/414 [00:00
<?, ?B/s]
```

```
In [16]: ## Preprocess dataset
max_length = get_max_length(model)
dataset = preprocess_dataset(tokenizer, max_length, seed, dataset)
```

```
Found max length: 4096
Preprocessing dataset...
Map: 0%|          | 0/15011 [00:00<?, ? examples/s]
Map: 0%|          | 0/15011 [00:00<?, ? examples/s]
Filter: 0%|          | 0/15011 [00:00<?, ? examples/s]
```

```
In [17]: def train(model, tokenizer, dataset, output_dir):
    # Apply preprocessing to the model to prepare it by
    # 1 - Enabling gradient checkpointing to reduce memory usage du
    model.gradient_checkpointing_enable()

    # 2 - Using the prepare_model_for_kbit_training method from PEF
    model = prepare_model_for_kbit_training(model)

    # Get lora module names
    modules = find_all_linear_names(model)

    # Create PEFT config for these modules and wrap the model to PE
    peft_config = create_peft_config(modules)
    model = get_peft_model(model, peft_config)

    # Print information about the percentage of trainable parameter
    print_trainable_parameters(model)

    # Training parameters
    trainer = Trainer(
        model=model,
        train_dataset=dataset,
        args=TrainingArguments(
            per_device_train_batch_size=1,
            gradient_accumulation_steps=4,
            warmup_steps=2,
            max_steps=15,
            learning_rate=2e-4,
            fp16=True,
            logging_steps=1,
            output_dir="outputs",
            optim="paged_adamw_8bit",
        ),
        data_collator=DataCollatorForLanguageModeling(tokenizer, ml

    )

    model.config.use_cache = False # re-enable for inference to sp

    ### SOURCE https://github.com/artidoro/qlora/blob/main/qlora.py
    # Verifying the datatypes before training

    dtypes = {}
    for _, p in model.named_parameters():
        dtype = p.dtype
```

```

dtype = p.dtype
if dtype not in dtypes: dtypes[dtype] = 0
dtypes[dtype] += p.numel()
total = 0
for k, v in dtypes.items(): total += v
for k, v in dtypes.items():
    print(k, v, v/total)

do_train = True

# Launch training
print("Training...")

if do_train:
    train_result = trainer.train()
    metrics = train_result.metrics
    trainer.log_metrics("train", metrics)
    trainer.save_metrics("train", metrics)
    trainer.save_state()
    print(metrics)

###

# Saving model
print("Saving last checkpoint of the model...")
os.makedirs(output_dir, exist_ok=True)
trainer.model.save_pretrained(output_dir)

# Free memory for merging weights
del model
del trainer
torch.cuda.empty_cache()

output_dir = "results/llama2/final_checkpoint"
train(model, tokenizer, dataset, output_dir)

```

all params: 3,540,389,888 || trainable params: 39,976,960 || trainable
 e%: 1.1291682911958425
 torch.float32 302387200 0.08541070604255438
 torch.uint8 3238002688 0.9145892939574456
 Training...

You're using a LlamaTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.

[15/15 00:38, Epoch 0/1]

Step	Training Loss
1	2.544200
2	1.918400
3	1.645800
4	1.603600
5	1.578300
6	1.278700
7	1.326800
8	1.325600
9	1.312500

10	1.231300
11	1.102100
12	0.807400
13	1.138600
14	1.413800